

Specifications of Custom Protocol

Connection Oriented

- The protocol is connection oriented because it requires setting up a handshake before data can be sent and received. Note that the sender keeps sending start signals until the sender sends one back, starting the connection.
- A connection is also ended when an end signal is sent to the receiver from the sender. Note that the sender will keep sending end signals until the receiver sends one back, confirming the end of connection.
- When a connection starts, it is agreed that the first expected packet will be numbered 0, and then iterated up to the MAX_PACKET_NUMBER and then loop back to 0.

Pipelining

- The protocol implements a form of selected repeat. When a sent packet is not ACK'ed within a timeout duration, the sender will resend the packet.
- Additionally, multiple packets can be waiting for ACK at a given time. This means that the sender is not blocked waiting for an ACK.

Flow Control

- The protocol uses a sliding window. The window does not progress until the oldest packet is ACK'ed. This ensures that the receiver is not overwhelmed.

Congestion Control

- The protocol has a limit to how many packets can be un-ACK'ed at a time. This ensures that the network is not overwhelmed.

Reliability

- The protocol is reliable because packets are resent if they are not ACK'ed. This prevents data loss.
- Checksum: The protocol uses a CRC32 checksum to ensure data integrity by verifying that the packet's checksum matches the computed value when it's received.
- Byte encoding: Packets follow the format [1 byte packetType] + [1 byte seqNum] + [4 bytes checksum] + [data]

Receiver Specifications

- The receiver ensures that messages are printed in order. It saves out of order messages for later when it will print them.

Testing Procedures

Simulating Loss

- To simulate packet errors, the computeChecksum function intentionally calculates an incorrect checksum 10% of the time

End to End testing

Handshakes

```
[sysri@bcd07443da56 mp2 % python sender.py localhost 8080  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Sending handshake  
Handshaked confirmed  
Enter message (or "quit" to exit): |
```

```
[sysri@bcd07443da56 mp2 % python reciever.py  
Handhook, sending back  
|
```

Handshakes are continuously sent until one is received back

Basic Sender and Receiver

```
[sysri@bcd07443da56 mp2 % python sender.py localhost 8080  
Sending handshake  
Handshaked confirmed  
Enter message (or "quit" to exit): hello  
Enter message (or "quit" to exit): world  
Enter message (or "quit" to exit): |
```

```
[sysri@bcd07443da56 mp2 % python reciever.py  
Handshook, sending back  
recieved: hello  
recieved: world  
|
```

Sender and Receiver both work

ACK's

```
[sysri@bcd07443da56 mp2 % python reciever.py  
Handshook, sending back  
recieved: hello  
recieved: world  
|
```

```
[sysri@bcd07443da56 mp2 % python sender.py localhost 8080  
Sending handshake  
Handshaked confirmed  
Enter message (or "quit" to exit): hello  
Enter message (or "quit" to exit): Got ACK for 0  
world  
Enter message (or "quit" to exit): Got ACK for 1
```

Receiver sends back an ACK to sender

Random Loss

```
[sysri@bcd07443da56 mp2 % python reciever.py  
Handhook, sending back  
Got corrupted packet  
Got corrupted packet  
received packet out of order: 2
```

Receiver gets packet 0 and packet 1 as corrupted, and then packet 2 works. This is from when CHAOS (randomization amount) is set to 0.5 (50% chance to corrupt).

Selective Repeat and Ordered Output

```
Sender closed.  
sysri@bcd07443da56 mp2 % python sender.py localhost 8080  
Sending handshake  
Handshaked confirmed  
Enter message (or "quit" to exit): message 1  
Enter message (or "quit" to exit): message 2  
Enter message (or "quit" to exit): Got ACK for 0  
messageGot ACK for 1  
^R  
message 3  
Enter message (or "quit" to exit): message 4  
Enter message (or "quit" to exit): Got ACK for 2  
Got ACK for 3  
|
```

```

sysri@bcd07443da56 mp2 % python reciever.py
Handhook, sending back
Got corrupted packet
Got corrupted packet
recieved: message 1
Got corrupted packet
Got corrupted packet
recieved: message 2
Got corrupted packet
Got corrupted packet
recieved: message 3
Got corrupted packet
recieved: message 4

```

It outputs the packets in order and saves out of order packets in a buffer. Sender continuously resends packets that have not been ACK'd.

End Connection

```

sysri@bcd07443da56 mp2 % python reciever.py
Handhook, sending back
Got corrupted packet
Got corrupted packet
Receiver closed.

```

```

sysri@bcd07443da56 mp2 % python sender.py localhost 8080
Sending handshake
Handshaked confirmed
Enter message (or "quit" to exit): quit
Sender closed.

```

Connection is successfully closed after sender sends an end signal and receiver confirms back to sender. If the receiver doesn't confirm, the sender will keep trying to end connection.

Network Captures

Connection-Oriented Protocol

1	0.000000	127.0.0.1	127.0.0.1	UDP	38 56126 → 8080 Len=6
2	0.000738	127.0.0.1	127.0.0.1	UDP	38 56127 → 1024 Len=6

The protocol operates in a connection-oriented manner, as demonstrated by the consistent flow of packets between the sender and receiver. Even though the capture does not explicitly show a traditional handshake, our own protocol ensures reliable message delivery with proper

acknowledgment from the receiver. This continuous communication without loss or out-of-order packets points toward connection-oriented characteristics. It is likely that the sender awaits acknowledgment before sending further data, ensuring proper session management.

Acceptable Performance

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	127.0.0.1	127.0.0.1	UDP	38	56126 → 8080 Len=6
2	0.000738	127.0.0.1	127.0.0.1	UDP	38	56127 → 1024 Len=6
4	4.288789	127.0.0.1	127.0.0.1	UDP	43	56126 → 8080 Len=11
5	4.289834	127.0.0.1	127.0.0.1	UDP	38	56127 → 1024 Len=6
8	9.209645	127.0.0.1	127.0.0.1	UDP	43	56126 → 8080 Len=11
9	9.210706	127.0.0.1	127.0.0.1	UDP	38	56127 → 1024 Len=6
11	13.526619	127.0.0.1	127.0.0.1	UDP	43	56126 → 8080 Len=11
12	13.527713	127.0.0.1	127.0.0.1	UDP	38	56127 → 1024 Len=6
13	23.369746	127.0.0.1	127.0.0.1	UDP	43	56126 → 8080 Len=11
14	23.370805	127.0.0.1	127.0.0.1	UDP	38	56127 → 1024 Len=6
15	27.719722	127.0.0.1	127.0.0.1	UDP	43	56126 → 8080 Len=11
16	27.720841	127.0.0.1	127.0.0.1	UDP	38	56127 → 1024 Len=6
17	37.497106	127.0.0.1	127.0.0.1	UDP	38	56126 → 8080 Len=6
18	37.499656	127.0.0.1	127.0.0.1	UDP	38	56127 → 1024 Len=6
19	37.500811	127.0.0.1	127.0.0.1	UDP	38	56127 → 1024 Len=6
32	38.509926	127.0.0.1	127.0.0.1	UDP	38	56126 → 8080 Len=6

The protocol exhibits acceptable performance with round-trip times between the sender and receiver remaining within a reasonable range. For instance, the delay between packets is typically under 1 second, with some slight variations (e.g., a 4-second gap between some data packets), which could be due to network conditions or receiver processing time. There are no significant delays or retransmissions observed, demonstrating that the protocol operates efficiently under normal conditions.

Flow Control and Congestion Control

No.	Time	Source	Destination	Protocol	Length	Info
72	24.133321	127.0.0.1	127.0.0.1	UDP	38	54267 → 8080 Len=6
73	24.134085	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
74	30.941411	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
75	30.942473	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
76	31.888139	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
77	32.421947	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
78	32.423042	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
79	32.902895	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
80	32.912333	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
81	32.913435	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
82	33.351406	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
83	33.352460	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
84	33.917085	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
85	33.919543	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
88	43.360499	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
89	43.362746	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
92	53.372618	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
93	53.374834	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
94	66.109686	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
95	66.110771	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
99	88.513631	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
100	88.514734	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
105	139.404496	127.0.0.1	127.0.0.1	UDP	43	54267 → 8080 Len=11
106	139.405669	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
148	168.964697	127.0.0.1	127.0.0.1	UDP	38	54267 → 8080 Len=6
149	168.967342	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
150	168.969891	127.0.0.1	127.0.0.1	UDP	38	54268 → 1024 Len=6
151	169.980680	127.0.0.1	127.0.0.1	UDP	38	54267 → 8080 Len=6

Based on the packet capture analysis, we can infer that flow control is being actively employed in the protocol. The sender initially transmits bursts of data in rapid succession, but these bursts are followed by intentional pauses in transmission. The long delays between certain packets, especially the 10-50 second gaps, suggest that the sender is deliberately slowing down its transmission rate. This behavior prevents network congestion and ensures that the receiver is not overwhelmed by excessive data, demonstrating that flow control mechanisms are in place to manage the flow of data between sender and receiver effectively. The consistency of small packet sizes further supports the idea that the protocol is carefully controlling data transmission to maintain smooth and efficient communication.

The packet capture also shows evidence of congestion control in action, as indicated by the varying time intervals between packet bursts. Initially, the sender sends bursts of packets at a steady pace, but over time, the gaps between these bursts gradually increase (e.g., from around 33 seconds to 43 seconds between bursts). This suggests that the sender is adjusting its transmission rate to avoid overwhelming the network, a hallmark of congestion control mechanisms. Additionally, the size of the packets fluctuates (e.g., packets of 38 bytes and 43 bytes), further indicating dynamic adjustment of the sending rate based on the network's congestion state. The consistent alternating of small and larger bursts of packets, along with the increased delay between them, is consistent with a protocol that is actively managing congestion to maintain network stability and prevent packet loss.